

Case 1 - Wholesale Outreach & Automation

I wouldn't try to scale wholesale through high-volume generic outreach. For a premium brand, that could hurt perception. I would build a system that scales the research, organization, and follow-up process, while keeping the actual communication selective, relevant, and personal. The goal would be to make outreach more efficient without making it feel automated.

1. Identifying and organizing target retailers

I would start by defining Mumzers' ideal wholesale profile based on what is already working. Looking at the current brand and stockists, I would focus on stores like premium kidswear boutiques, elevated family lifestyle stores, luxury resort retail and/or premium gifting stores.

Assuming we should be much more interested in brand fit than simply volume, we'd build the target list using three sourcing methods:

Lookalike retailer sourcing: we'd analyze Mumzers' current stockists and identify other retailers with similar characteristics: similar price point, similar visual identity, similar customer, and similar product mix. If Mumzers already performs in with them, then that gives a very clear blueprint for where else the brand should be.

Adjacent brand sourcing: We'd look at other premium children's or family lifestyle brands that sit in a similar aesthetic and pricing, then review their stockists, "where to buy" pages, Instagram tags, and retailer mentions. That's a great way to build a qualified retailer universe because those stores have already shown they understand premium kids products and visual merchandising.

Market-led sourcing by geography/channel: We'd split prospects by market and retailer type, for example:

- UK and Europe premium independents, concept stores, resorts, and luxury family boutiques
- US premium department stores, better specialty stores, and high-end online retail
- travel, resort, and lifestyle-driven stores in markets where Mumzers' bold aesthetic naturally fits

Once identified, we'd put every prospect into a structured CRM with fields such as: store name, website, geography, store type, number of locations, founder name, email, Instagram, current brands stocked, aesthetic fit, price positioning, why they are a fit, outreach status, last touch date, next step, and priority score.

I would also score retailers into clear tiers (ex: A- high-fit, high-value, strategically important accounts; B- strong-fit independents and online stores with good potential; C- possible fit, but not immediate priority).

2. Outreach and follow-ups

I would structure outreach in a way that feels consistent internally, but personalised externally.

A- Highly tailored outreach. Referencing something specific about the store: their customer, their market, or the brands they already carry. The message should feel like Mumzers chose them, not like they were one of 1000000000 stores added to a sequence.

B- Semi-personalized outreach. The core message can stay consistent, but parts of it should change depending on the retailer type, geography, and why they are a fit. That keeps quality high while allowing scale.

C- I'd keep them in a lighter nurture flow until there is a stronger reason to prioritize them, such as a new collection, a stockist win in their market, or a seasonal angle that makes the product more relevant.

Example outreach sequence (of course having data on what works already will help us curate these):

Email 1: short intro, why we think the store is a fit, brief brand positioning, and a clear call to action.

Email 2: follow-up with added value, for example hero products, best-seller price points or existing stockists.

Email 3: another touchpoint with a different angle, such as styling versatility or seasonality.

Email 4: more personal follow-up, potentially in a founder voice, asking whether timing is the issue or whether someone else is the correct buyer.

Email 5: polite close-the-loop email that keeps the door open without dragging the thread forever.

Every follow-up should bring something new. For a premium brand, repeatedly sending "hi, have you noticed us?" weakens the impression. Instead, each touch should add relevance, proof, or context.

->Since Mumzers already has an active social presence, it makes sense to use Instagram to build familiarity with carefully chosen retailers before or alongside email outreach. Always carefully because luxury brands can look cheap very fast if they overdoo direct messaging.

3. Tools and automation processes

I'd keep the tech stack practical and clean. We could use CRMs like HubSpot or Airtable to manage the pipeline from prospecting to order.

We could also use a simple “internal operations machine” (airtable/notion) that would include: current line sheets, lookbooks, wholesale decks, stockist proof, brand story assets, etc.

The automation itself should handle the repetitive work, creating and tagging prospect records, assigning them to a tier, trigger the right outreach flow, followup reminders, etc

Simply: automating the workflow, not the relationship.

4. Maintaining quality

If the brand is premium, the outreach process has to reflect that. I would rather reach 100 highly relevant retailers properly than 1000 loosely relevant ones badly. Mumzers already has proof that it belongs in curated, high taste environments, so I would not damage the brand trying to chase every store with a children's section.

Top accounts should feel researched and intentionally selected. If Mumzers is presenting itself as design-led and luxury-leaning, the wholesale deck, line sheet, imagery, and sample presentation all need to feel polished and consistent with the site and product experience. The B2B materials and communication need to match the rest of the brand as well.

Wholesale success should not just be measured by the number of doors opened, but by the quality of the doors.

How

P1: Building our environment Mapping current stockists, defining the ideal retailer profile, sourcing lookalike prospects, and organizing them into a CRM with scoring and tiers.

P2: Building the outreach system Creating outreach templates by tier, defining follow-up cadences, organizing sales assets, and setting clear pipeline stages.

P3: Measuring and refining: Tracking not just opens and replies, but positive reply rate, meetings booked, conversion to account, first order value, and reorder rate by geography and retailer type. Then doubling down on the segments that perform the best.

My overall approach would be to turn Mumzers' current wholesale process from something manual into a system that is structured, repeatable, and scalable, while still feeling premium.

In summary, systemizing the back end, while keeping the front end curated, selective, and human.

Case 2 — Increasing Consistency in Direct-to-Consumer Sales

I would treat it as a consistency problem inside the DTC engine. If customers already place strong orders when they do buy, good. It suggests the product, offer, or basket potential is already there. The question is why that demand is not showing up more predictably.

I'd try to understand where the inconsistency is happening, then build a more reliable system across traffic, conversion, and retention. The goal would be to reduce randomness and create a steadier sales rhythm rather than relying on occasional spikes.

1. Data

I would start by breaking the DTC performance down into the core components: traffic, conversion rate, average order value, and repeat purchase behavior.

I'd look into 6 to 12 months of Shopify data to understand: daily and weekly revenue patterns, session volume trends, conversion rate trends, average order value trends, sales by channel top products.

I'm essentially looking for the answer to the question: are quiet weeks happening because fewer people are arriving, because fewer visitors are converting, or because repeat customers are not coming back often enough?

I would then want to know which sources are driving the strongest customers and which ones are creating volatility. For example, are strong weeks tied to ad spendings, emails sent, influencer/partner activity, posting, or seasonal moments?

If so: then part of the problem may simply be that the business is too dependent on irregular traffic drivers instead of having a more stable acquisition and retention system.

I'd also review on-site behavior and funnel performance. That means looking at: landing pages' traffic; product page conversion rates, add-to-cart rates, cart-to-checkout rate, checkout completion rate, multiple device performance and page friction (if any).

Since order values are already strong, we should also pay attention to what customers are buying together. If there are products that consistently drive stronger baskets, those should probably be analyzed further.

Finally, I would review retention and lifecycle performance to understand: repeat purchase rate, time between first and second order, email capture rate, abandoned cart performance and post-purchase performance.

If new customers buy well but there is no structured system bringing them back, that will naturally create inconsistency over time.

Most of these insights can be checked in shopify directly. If not all, google search console and others should help.

2. Possible strategies for consistent traffic and sales

These would depend on the inconsistencies, but the goal would be the same: predictability.

First priority: creating more consistent traffic sources. If paid acquisition is already part of the mix, I would suggest a steadier always-on structure rather than bursts of activity (assumption). Keeping a reliable prospecting layer live, a clear retargeting layer, and a regular creative testing rhythm so traffic is not overly dependent on a small number of ad winners or short campaign windows.

On the other hand: I would strengthen owned channels to reduce volatility. Email/SMS should not just be occasional promotional tools. As referred in case 1, they should be part of a system. We should have a strong abandoned cart flow, post-purchase flow, and winback flow in place, supported by a regular campaign calendar. That helps create a base level of recurring traffic and revenue that is not fully dependent on spending money.

On the site itself: I would test ways to make conversion more stable. If customers already tend to place strong first orders, the next improvements could include: reviewing product page messaging, stronger social proof and trust signals, frequently bought together recommendations, new landing pages built around hero products, etc.

3. How

In the first 30 days, I would focus on diagnosis. We could create a weekly dashboard pulling data from Shopify, GA4, the email platform, and paid channels. I'd keep it focused on the core DTC components: traffic, conversion and retention. Structurally, I'd break it into headline KPIs, channel performance, funnel and landing-page performance, and product/retention performance.

For a first version I'd use sheets (or Looker Studio, I haven't used it but seems like a good option) so we can use it quickly, then automate it once the structure proves useful. The purpose would not be reporting just because, but creating a weekly decision-making system around what changed, why it changed, and what to do next. That would create a shared view of performance and make it much easier to spot patterns.

Then I'd audit the funnel end to end. That includes the Shopify storefront, product pages, cart and checkout experience, mobile experience, lifecycle flows, tracking setup, and channel performance. I would be looking for obvious leaks or missed opportunities that can be fixed quickly.

Finally, I'd check and try to improve the retention foundation. If flows are weak or missing, we should prioritize them (welcome flow, abandoned cart flow, post-purchase flow, etc).

After the first 30 days, I would move on to experimentation.

I'd build a backlog (airtable, notion, sheets) with hypotheses, owners, timelines, and success metrics. I would want every experiment to answer real valuable questions like "are we attracting the wrong traffic" or "what's turning (or not) interested people into convinced buyers".

It's also important to have a weekly operating rhythm: the team should be consistently reviewing the same components: traffic quality, funnel performance, conversion tests, lifecycle performance, and next experiments. That matters because consistency in sales usually comes from consistency in execution.

Some notes

An exercise like this would require more research, planning, and a better view of available resources and budget before execution. Some of my assumptions could be wrong, for example, because they are based on previous experiences. Even so, I think this is a good example of where I'd start.

I want to be clear that I used AI for research, to help with certain decisions, to brainstorm possibilities and would also use it usually before picking any tools (nowadays there's a new crazy magic tool coming out every day). I usually do not work/present work, or suggest solutions based on gut feeling. I prefer to work with facts, patterns, and available data while, when possible/needed, I can also be a very creative person.

Regarding tools: choosing a setup would depend on budget. Some tools will be stronger than others depending on the specific use case, so I would first assess the actual workflow and where the real bottlenecks are before deciding on the tools to use. Free or lower-cost options are possible too, but time is a clear tradeoff, since they often require custom building.

A few small things I noticed while on the site while working:

- the "20% off" message could probably benefit from a slight font size increase
- in checkout, the "you may also like" slider is not very user-friendly to click through
- when browsing on Apple devices, I was not prompted to use Apple Pay on checkout
- on the Swimzers page, there is a "no products" message showing, i suppose it's not intentional